


COMP3251

Lecture 9: Dijkstra Algorithm
(Chapter 4.3, 4.4, and 4.5)

Edges with Lengths

We now consider weighted graph, i.e., every edge (u,v) is associated with a length $L(u,v)$.




Single-Source Shortest Paths Problem (for Weighted Graphs)

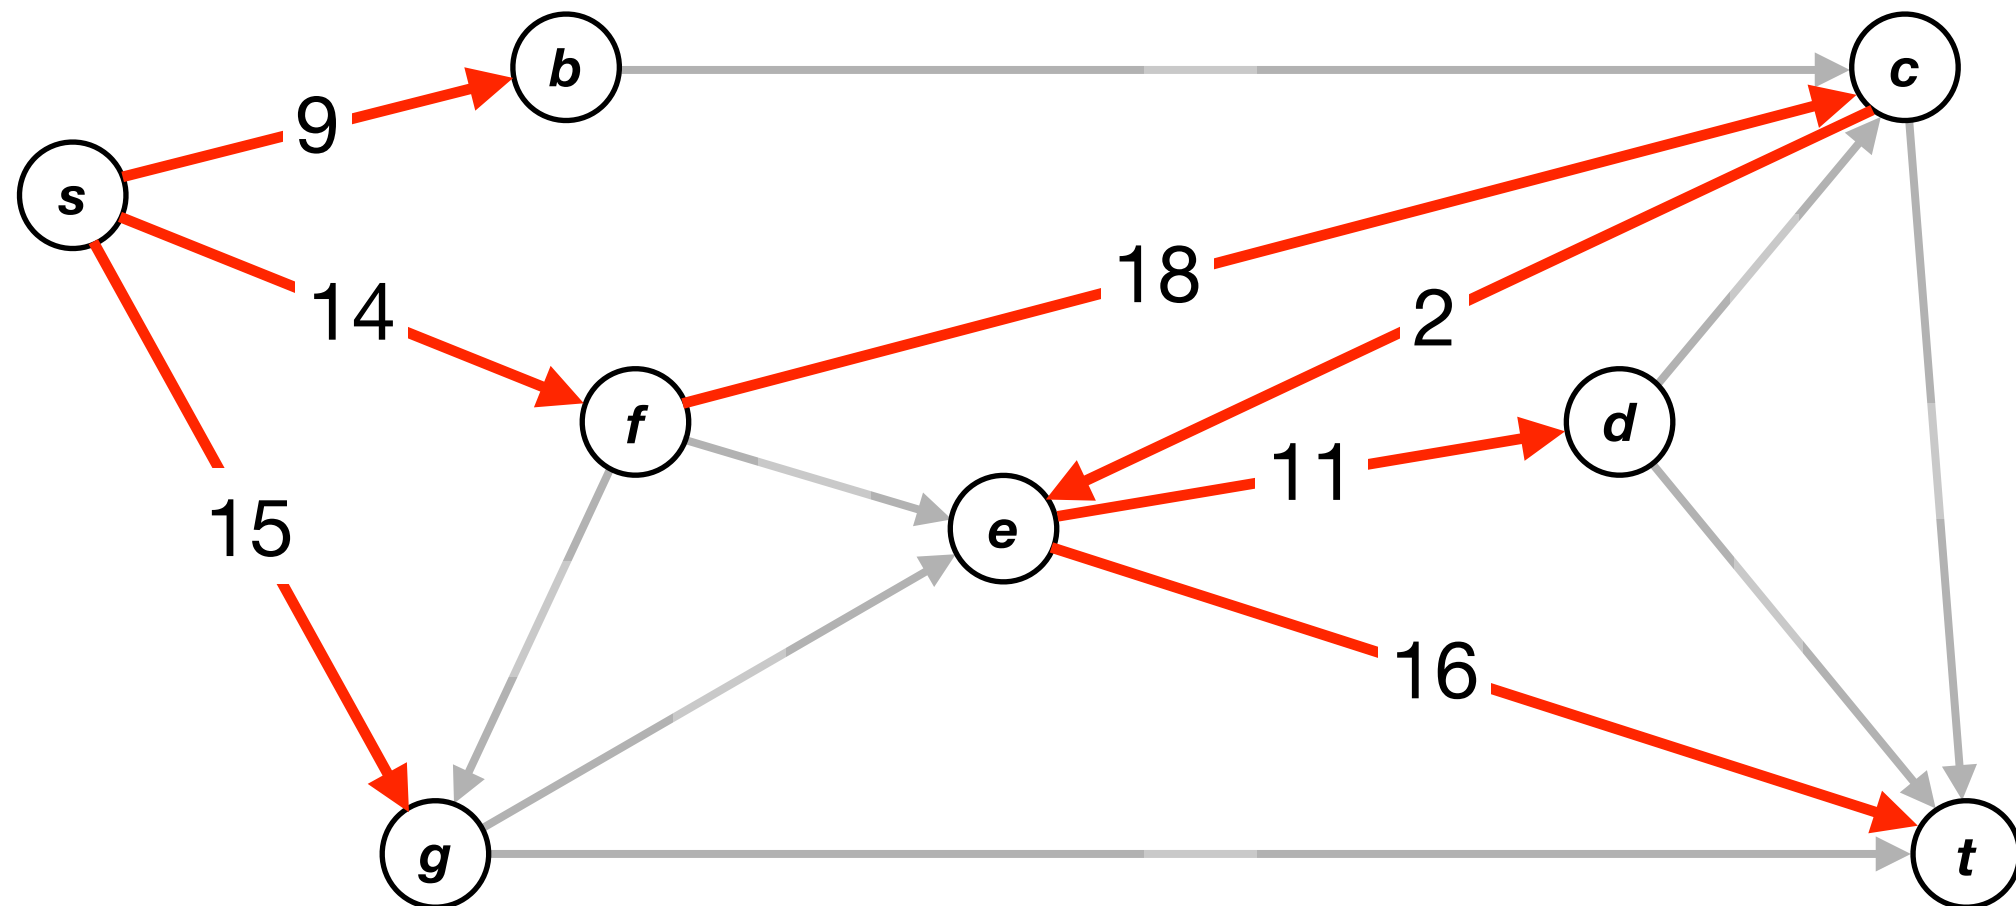
Given a weighted directed graph and a source vertex $s \in V$, find, for every vertex $v \in V$, the length of the shortest path from the source s to v .

Definition: length of a path P is equal to the sum of the length of the edges on P .

An Example




An Example



Note that this is a tree, and for any vertex v , the tree path from s to v is the shortest path s to v . Such a tree is called the Shortest Path Tree (SPT).

The BFS tree is a Shortest Path Tree

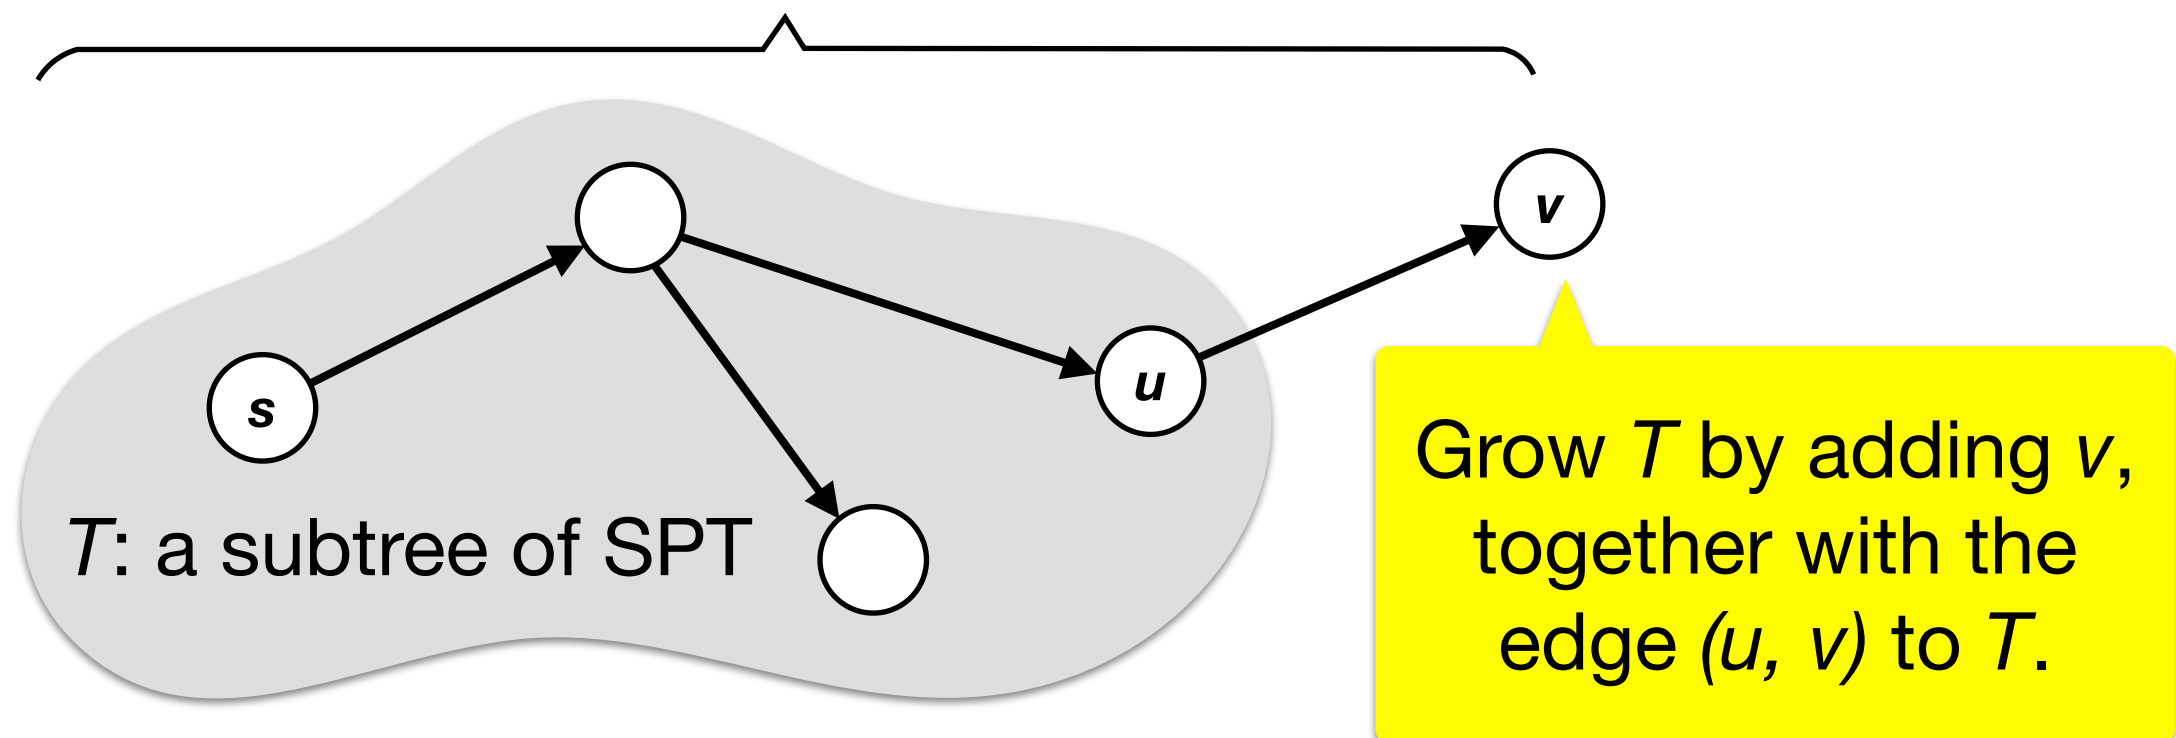
- If all edges have length 1, then BFS solves the shortest path problem and the BFS tree coincides with the SPT tree.
- An alternative view of BFS:
 - Starting from the smallest subtree of BFS tree that contains only the starting vertex s .
 - Add vertices and the corresponding tree edges one by one in ascending order of their distance from s .



General Idea

- Starting from the smallest subtree of SPT that contains only s .
- Iteratively attached a “correct” edge to the subtree such that the larger tree is still a subtree of SPT.
- When we include all vertices in the subtree, it is the SPT.

A larger T which is still a subtree of SPT




What is the correct vertex/edge?

Let T be a subtree of the SPT. Here are a few **definitions**:

- A T -path from s to v is a path from s to v that goes through only vertices in T .
- The T -dist(v) is the length of shortest T -path from s to v .
- Note that for any vertex v , T -dist(v) \geq dist(v), the length of the shortest path from s to v .
- For every $v \in T$, T -dist(u) = dist(u).

Example:




What is the correct vertex/edge?

Observation: A shortest T -path from s to v , must be consists of a shortest T -path from s to some vertex $u \in T$ and edge (u, v) .

- To determine $T\text{-dist}(v)$, we check all vertices $u \in T$ such that $(u, v) \in E$, and the corresponding shortest T -path from s to u .
- Then, determine their total lengths and return the minimum.

Example:



$T\text{-dist}(a) + L(a, v) = 35$,
 $T\text{-dist}(u) + L(u, v) = 34$,
 $T\text{-dist}(b) + L(b, v) = 38$,
so $T\text{-dist}(v) = 34$.

The Algorithm (Informal)


Starting with $T = \{s\}$ and repeatedly do the following to enlarge T until T includes all vertices:

- Find the vertex $v \notin T$ with the smallest $T\text{-dist}(v)$, add it to T ;
- Update the $T\text{-dist}(x)$ for all vertices $x \notin T$.

What next?

- Why is the above algorithm correct?
- What is its running time?

Correctness




Suppose:

- T is a subtree of SPT.
- v has the smallest T -dist among vertices that are not in T , the the path leaves T through a vertex $u \in T$.


Want to show:

- This path from s to v is indeed the shortest path.

Correctness (cont'd)




Correctness (cont'd)



By our choice of v , for all $x \notin T$,
 $T\text{-dist}(v) \leq T\text{-dist}(x)$

Correctness (cont'd)

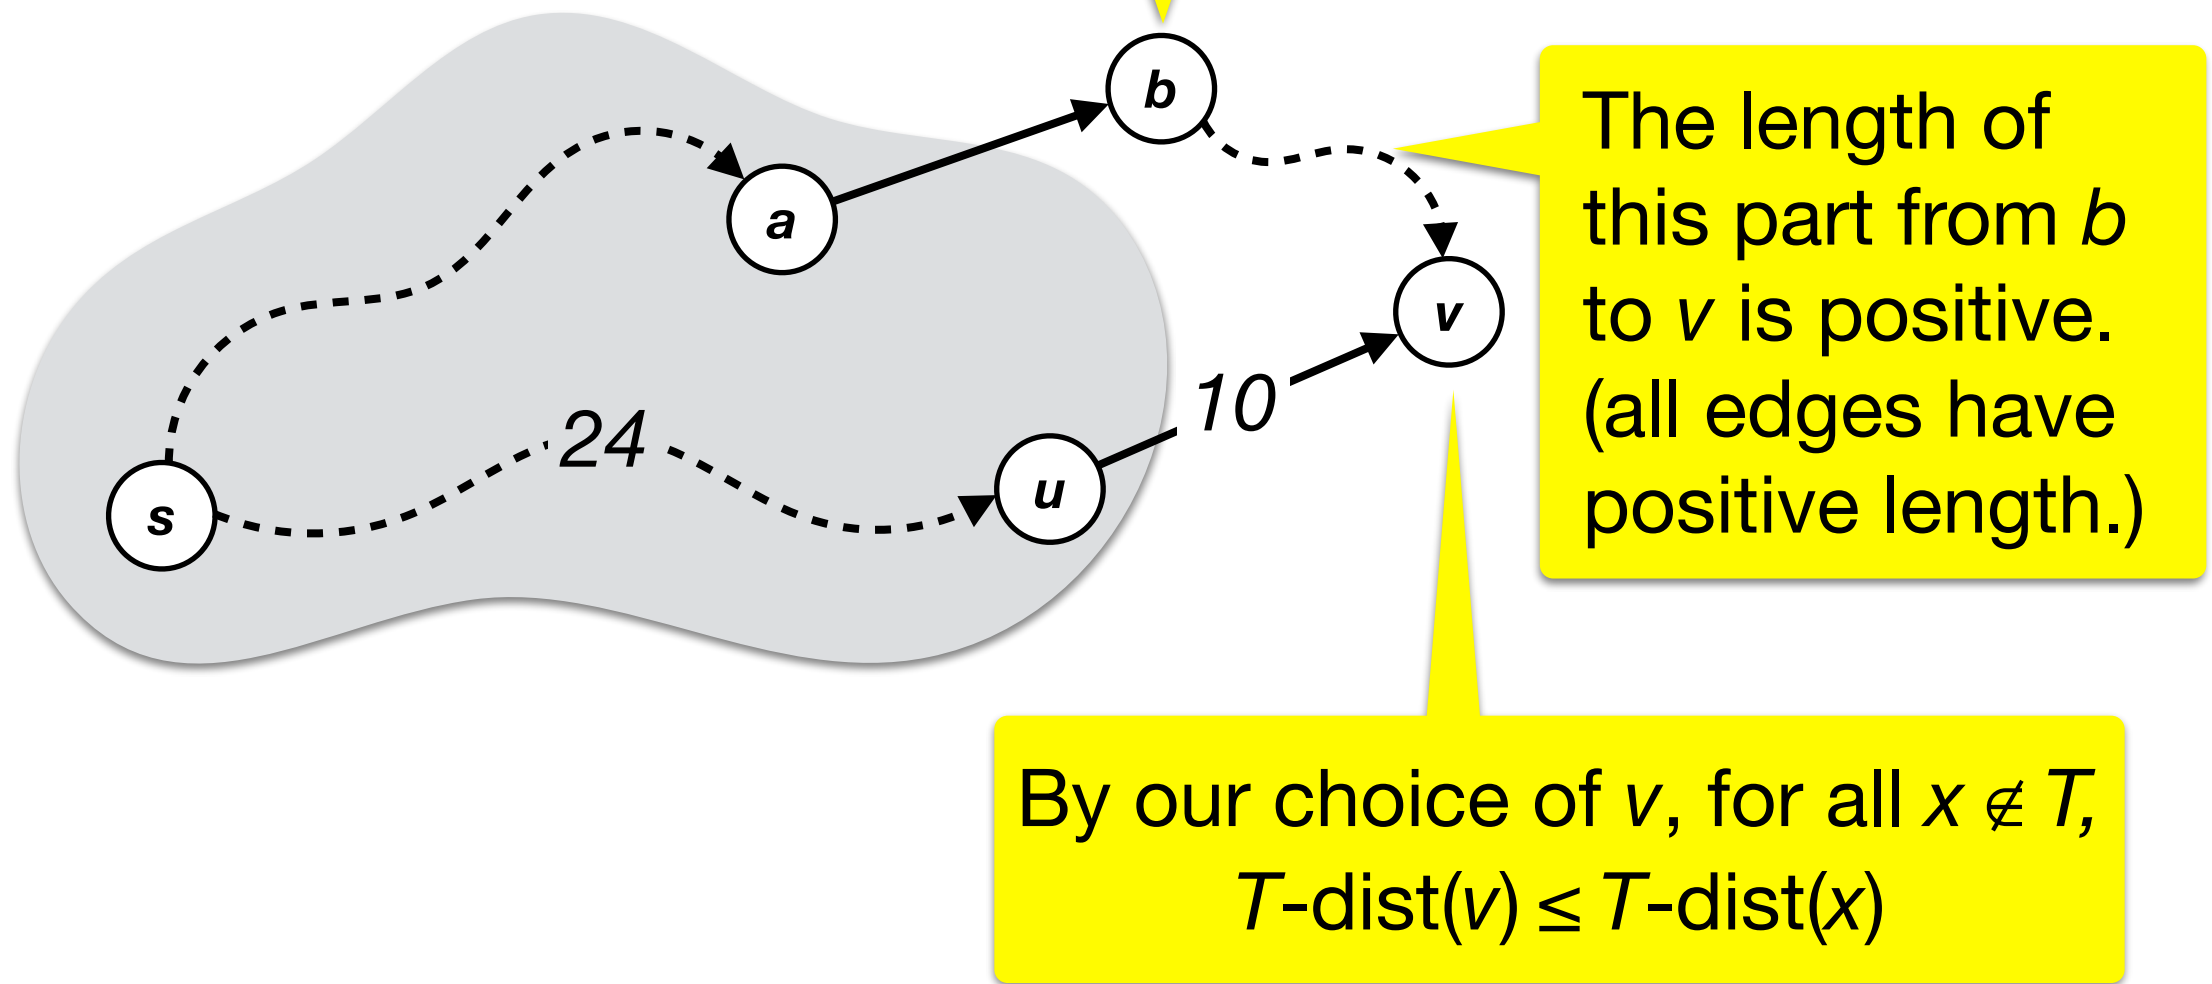
Since going through u is the shortest T -path from s to t , any alternative path from s to v must go through some vertex $b \notin T$, and $T\text{-dist}(b) \geq T\text{-dist}(v)$.



By our choice of v , for all $x \notin T$,
 $T\text{-dist}(v) \leq T\text{-dist}(x)$


Correctness (cont'd)

Since going through u is the shortest T -path from s to t , any alternative path from s to v must go through some vertex $b \notin T$, and $T\text{-dist}(b) \geq T\text{-dist}(v)$.




Correctness (cont'd)

Conclusion: The length of any path P from s to v is no smaller than $T\text{-dist}(v)$ because $\text{length}(P) \geq T\text{-dist}(b) \geq T\text{-dist}(v)$. Hence, the T -path from s to v through u is the shortest path.



Sample Run


Highlighted numbers denote the length of the shortest path from s to a vertex



$$T = \{s\}$$

Sample Run


Shaded numbers denote the current T -distance of a vertex



$$T = \{s\}$$

Sample Run


Add the vertex $v \notin T$ with the smallest T -dist




$$T = \{s, a\}$$

Sample Run

Update $T\text{-dist}(x)$ for all $x \notin T$




Sample Run




$$T = \{s, a, d\}$$

Sample Run




$$T = \{s, a, d\}$$

Sample Run




$$T = \{s, a, d, e\}$$

Sample Run




$$T = \{s, a, d, e\}$$

Sample Run




$$T = \{s, a, d, e, b\}$$

Sample Run




$$T = \{s, a, d, e, b, c\}$$

How to update $T\text{-dist}(x)$?




- Recall $T\text{-dist}(x) = \min \{ T\text{-dist}(a) + L(a, x), T\text{-dist}(b) + L(b, x) \}$
- Suppose the algorithm adds a new vertex v to T .
 - Let T' denote the new tree.
 - Let $T'\text{-dist}$ denote the updated tree-distance.

How to update $T\text{-dist}(x)$?




- If there is no edge (v, x) , then we have
$$T'\text{-dist}(x) = \min \{ T'\text{-dist}(a) + L(a, x) , T'\text{-dist}(b) + L(b, x) \}$$
- Adding v to T does not change the tree-distance of a and b :
 $T'\text{-dist}(a) = T\text{-dist}(a)$, $T'\text{-dist}(b) = T\text{-dist}(b)$. Hence, no change.

How to update $T\text{-dist}(x)$?



- If there is edge (u, x) , then $T'\text{-dist}(x)$ is equal to $\min \{ T'\text{-dist}(a) + L(a, x) , T'\text{-dist}(b) + L(b, x) , T'\text{-dist}(v) + L(v, x) \}$
- Note that $T\text{-dist}$ and $T'\text{-dist}$ are the same for v, a, b . Thus,
$$T'\text{-dist}(x) = \min \{ T\text{-dist}(x) , T\text{-dist}(v) + L(v, x) \}$$

How to update $T\text{-dist}(x)$?



Conclusion: After adding v to enlarge T , we only need to

- scan the adjacent list of v in order to update $T\text{-dist}$;
- for each vertex x in the adjacent list, update its tree-distance as the minimal of the old tree-distance and $T\text{-dist}(v) + L(v, x)$.

The Algorithm (DPV 4.1.1)

- 1) **initialize** $\text{dist}(s) = 0$, and $\text{dist}(x) = \infty$ for other vertices x ;
- 2) $V' = \{ \}$;
- 3) **while** $V' \neq V$ **do**
- 4) pick the node $v \notin V'$ with the smallest $\text{dist}(\cdot)$
- 5) add v to V'
- 6) **for** all edges $(v, x) \in E$:
- 7) **if** $\text{dist}(x) > \text{dist}(v) + L(v, x)$: update $\text{dist}(x) = \text{dist}(v) + L(v, x)$.

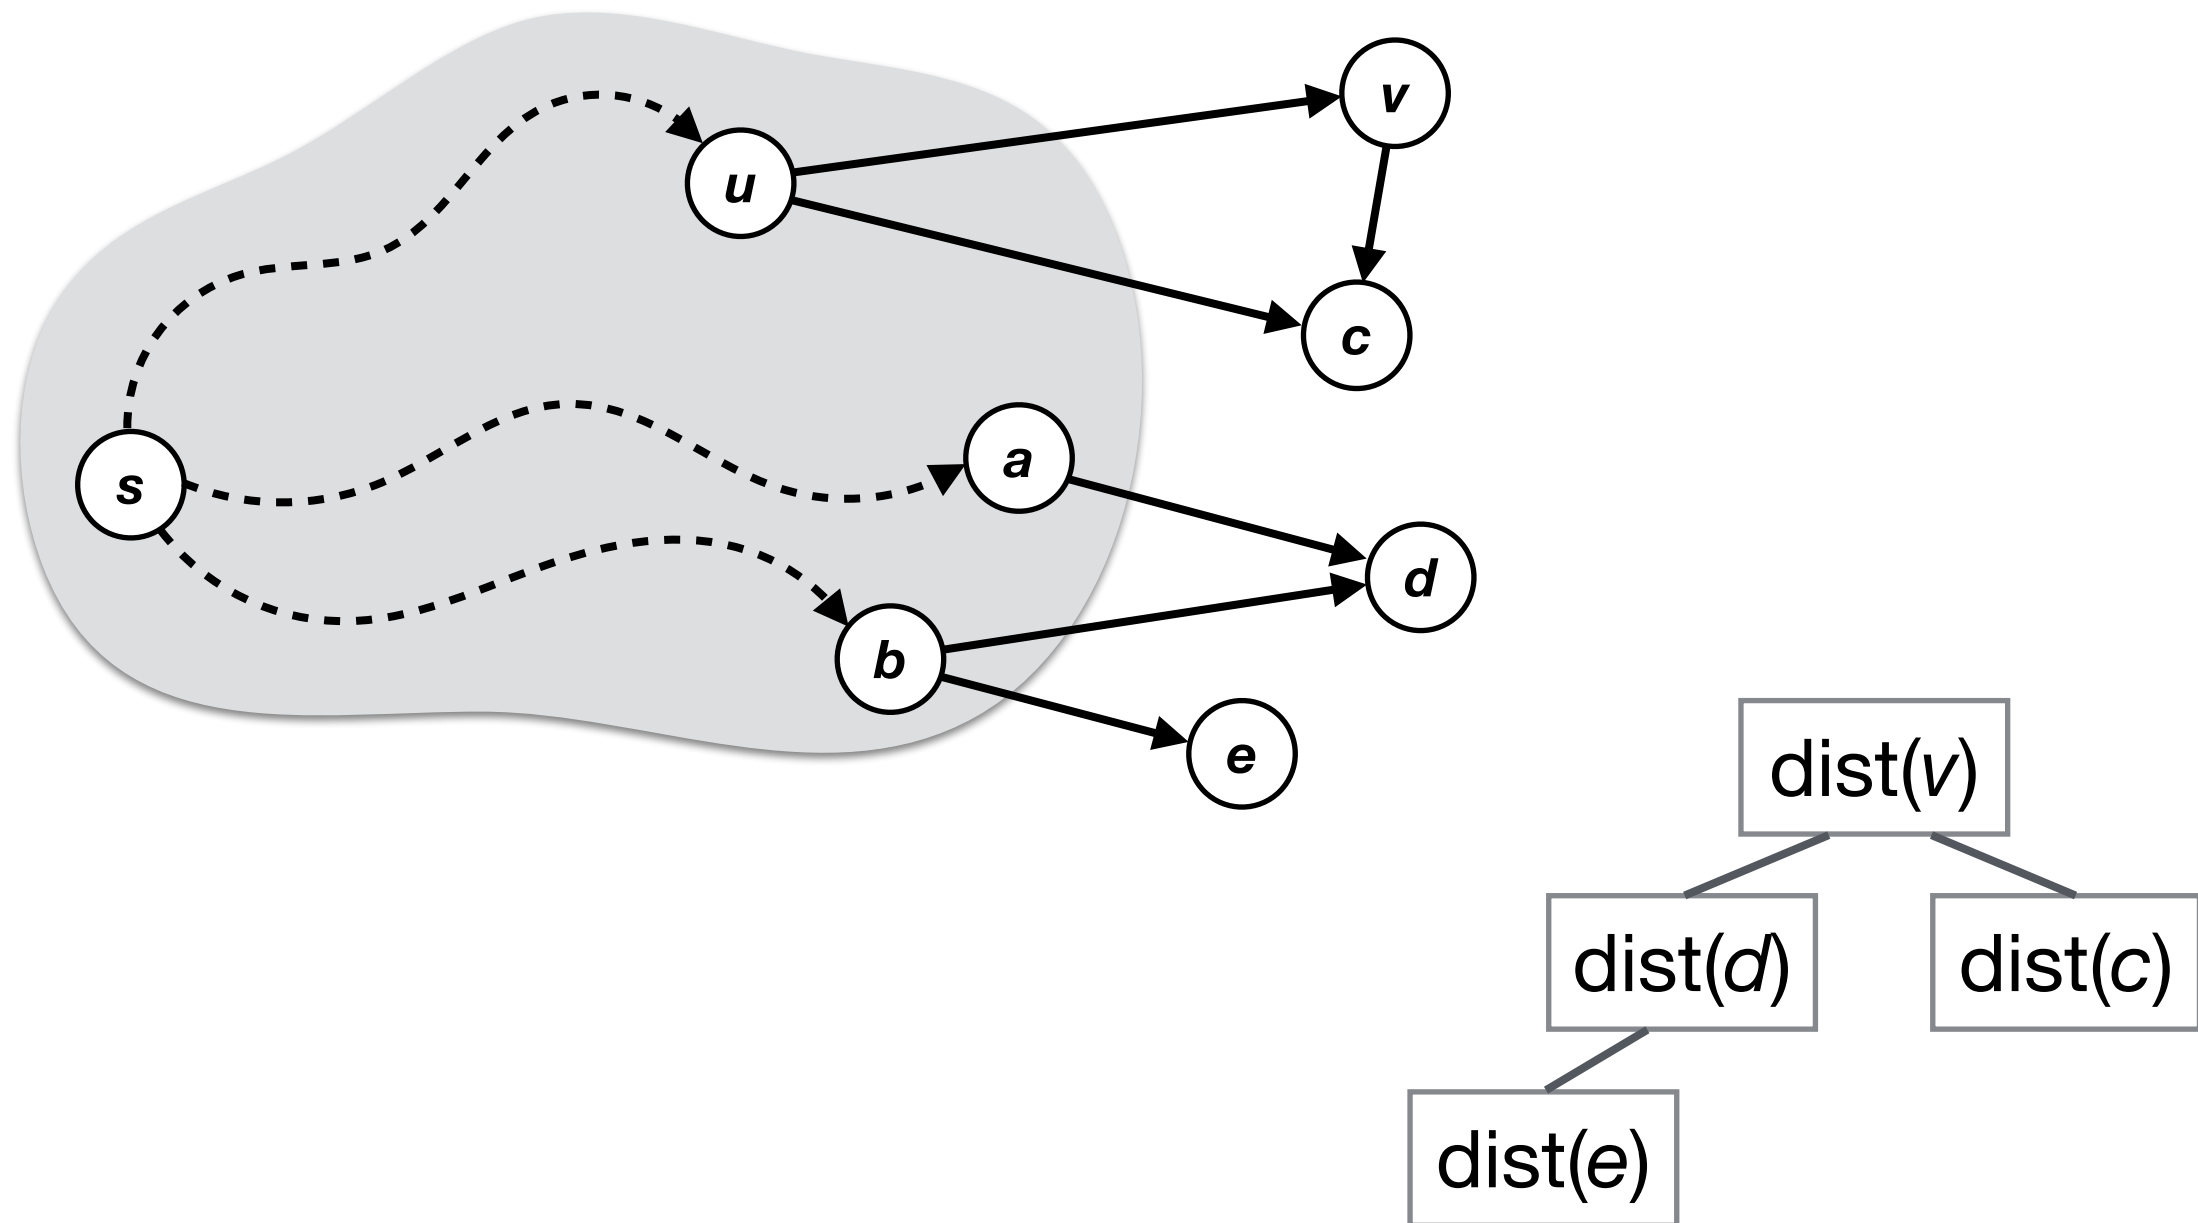
Running time:

- Steps 1, 2, 5 take $O(|V|)$ time. Steps 6, 7 takes $O(|E|)$ time.
- If we implement $\text{dist}(\cdot)$ as a simple array, step 4 takes $O(|V|)$ time to find the smallest $\text{dist}(\cdot)$, and since we execute step 4 $O(|V|)$ times, total time is $O(|V|^2)$.


Priority Queue Implementation

- It is usually more efficient to store $\text{dist}(\cdot)$ in a priority queue, e.g., implemented through binary min heap.
- Recap of basic properties of binary min heap:
 - Building an empty heap (**build-heap**): $O(1)$;
 - Finding the minimal number (**find-min**): $O(1)$;
 - Deleting a the minimal number (**delete-min**): $O(\log n)$.
 - Inserting a new number (**insert**): $O(\log n)$;


Priority Queue Implementation




Priority Queue Implementation




find-min




Priority Queue Implementation




delete(v)




Priority Queue Implementation




Update $\text{dist}[x]$ for every x s.t. there is edge (v, x) ; for example, $\text{dist}[c]$ may be decrease by some $\Delta > 0$.



Priority Queue Implementation




Remember, we may need to shift up the new value in order to restore heap order.



Running time for this implementation

- Initially, we set $\text{dist}(s) = 0$, and $\text{dist}(u) = \infty$ for all other vertices u . Then, we insert all n vertices to a heap: $O(|V| \log |V|)$ time.
- We need to delete the minimum value from the heap n times to add n vertices to the SPT: $O(|V| \log |V|)$ time.
- After adding each vertex to the SPT, we need to scan its adjacency list and for each vertex x in the list, we may need to update $\text{dist}(x)$ and shift up to restore heap order.
 - The height of the tree is $\log |V|$, so we need to shift up at most $\log |V|$ times per vertex in the adjacency list.
 - Total size of the adjacency lists of all vertices is $O(|E|)$.
 - In total, this takes $O(|E| \log |V|)$ time.
- In sum, the running time is $O((|E| + |V|) \log |V|)$.

Advance topic: d -heap implementation



	Heap	3-Heap	d -Heap
build-heap	$O(1)$	$O(1)$	$O(1)$
insert	$O(\log n)$	$O(\log_3 n)$	$O(\log_d n)$
find-min	$O(1)$	$O(1)$	$O(1)$
delete-min	$O(\log n)$	$O(3 \log_3 n)$	$O(d \log_d n)$

Advance topic: d -heap implementation

- Initially, we set $\text{dist}(s) = 0$, and $\text{dist}(u) = \infty$ for all other vertices u . Then, we insert all n vertices to a heap: $O(|V| \log_d |V|)$ time.
- We need to delete the minimum value from the heap n times to add n vertices to the SPT: $O(|V| d \log_d |V|)$ time.
- After adding each vertex to the SPT, we need to scan its adjacency list and for each vertex x in the list, we may need to update $\text{dist}(x)$ and shift up to restore heap order.
 - The height of the tree is $\log_d |V|$, so we need to shift up at most $\log |V|$ times per vertex in the adjacency list.
 - Total size of the adjacency lists of all vertices is $O(|E|)$.
 - In total, this takes $O(|E| \log_d |V|)$ time.
- Total $O((|E| + d |V|) \log_d |V|)$ time.
- Choosing $d = |E|/|V|$, the running time is $O(|E| \log_{|E|/|V|} |V|)$.