

COMP3251


Lecture 12: Kruskal's Algorithm
(Chapter 5.1)

Minimum Spanning Tree

Definition: Given a **connected undirected graph** $G = (V, E)$ in which every edge $e \in E$ is associated with a positive **weight** $w(e)$, a **minimum spanning tree (MST)** is a subset of edges $T \subseteq E$ s.t.

- (i) T forms a spanning tree; and
- (ii) the sum of edge weights of T is minimized.

Example:



Two Greedy Algorithms for MST

Prim's algorithm (last lecture)

- Start with some root node s and grow a tree T outward.
- At each step, add the minimum weight outgoing edge.
- This algorithm is almost the same as the Dijkstra's algorithm, except that we add the outgoing edge with the minimum weight, not the one with minimum T -distance.

Kruskal's algorithm (this lecture)


- Start with T being the empty forest.
- Consider edges in ascending order of cost; insert edge e in T unless doing so would create a cycle.

Kruskal's Algorithm (Chapter 5.1.3)

- 1) Start from an empty forest T .
- 2) **for** all edges e in ascending order of weights :
- 3) insert edge e in T unless doing so would create a cycle.

Kruskal's Algorithm (Chapter 5.1.3)


- 1) Start from an empty forest T .
- 2) **for** all edges e in ascending order of weights :
- 3) insert edge e in T unless doing so would create a cycle.



Initially, T contains no edges.

Kruskal's Algorithm (Chapter 5.1.3)


- 1) Start from an empty forest T .
- 2) **for** all edges e in ascending order of weights :
- 3) insert edge e in T unless doing so would create a cycle.



At step 1, the minimum weight edge we could add has weight 1.

Kruskal's Algorithm (Chapter 5.1.3)


- 1) Start from an empty forest T .
- 2) **for** all edges e in ascending order of weights :
- 3) insert edge e in T unless doing so would create a cycle.



At step 2, the minimum weight edge we could add has weight 3.

Kruskal's Algorithm (Chapter 5.1.3)


- 1) Start from an empty forest T .
- 2) **for** all edges e in ascending order of weights :
- 3) insert edge e in T unless doing so would create a cycle.



At step 3, the minimum weight edge we could add has weight 5.

Kruskal's Algorithm (Chapter 5.1.3)

- 1) Start from an empty forest T .
- 2) **for** all edges e in ascending order of weights :
- 3) insert edge e in T unless doing so would create a cycle.




At step 4, the edge with weight 6 creates a cycle, so we add the edge with weight 7 instead.

Correctness of Kruskal's Algorithm


Proof by picture:

- Recall that if e is the minimum weight outgoing edge of some subset of vertices S , then the MST must contain e .
- For every edge e we add in the sample run, we will explain the subset of vertices S for which e is the minimum weight outgoing edge of S , certifying e must be in the MST.




Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(h,g) is the minimum weight edge going out from $S = \{ h \}$.

Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(i,c) is the minimum weight edge going out from $S = \{ c \}$.

Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(g,f) is the minimum weight edge going out from $S = \{ f \}$.

Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(a,b) is the minimum weight edge going out from $S = \{ a \}$.


Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(c,f) is the minimum weight edge going out from $S = \{ c, i \}$.


Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(c,d) is the minimum weight edge going out from $S = \{ d \}$.


Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(h,i) cannot be added to the solution because doing so would create a cycle.


Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(b,c) is the minimum weight edge going out from $S = \{ a, b \}$.


Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(a,h) cannot be added to the solution because doing so would create a cycle.


Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



Correctness of Kruskal's Algorithm


(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



(d,e) is the minimum weight edge going out from $S = \{ e \}$.


Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



Correctness of Kruskal's Algorithm

(h,g)	1
(i,c)	2
(g,f)	3
(a,b)	4
(c,f)	5
(c,d)	6
(h,i)	7
(b,c)	8
(a,h)	9
(d,e)	10
(e,f)	11
(b,h)	12
(d,f)	14



We already have $|V|-1$ edges. So adding any of the remaining edges would create a cycle.

Correctness of Kruskal's Algorithm

Fact 1. The algorithm adds an edge only if it is in the MST.

Key question: Why we can always find a subset of vertices S such that e is the minimum weight outgoing edge from S ?

- Throughout the algorithm, we have a set of subtrees.
- We add an edge e only if it does not create any cycle.
- So the two endpoints of e cannot be in the same subtree. That is, e connects two different subtrees, say T_1 and T_2 .
- So e is an outgoing edge of T_1 . Further, the algorithm has not processed any outgoing edges of T_1 when we add e .
- Choosing $S = T_1$ suffices because by our choice of e , it must have the minimum edge weight among them.

Correctness of Kruskal's Algorithm

Fact 2. Each edge in the MST will be added by the algorithm.

Proof: Consider an edge e in the MST.

- Since the algorithm checks all edges before it stops, it must have checked edge e .
- Since the only edges added by the algorithm are those in the MST, e cannot create a cycle.
- So the algorithm would have added edge e to the solution.

Implementing Kruskal's Algorithm (A Data Structure for Disjoint Sets)

Implementing Kruskal's Algorithm


- 1) Start from an empty forest T .
- 2) **for** all edges e in ascending order of weights :
- 3) insert edge e in T unless doing so would create a cycle.

- Note that $\log |E| = O(\log |V|)$ because $|E| = O(|V|^2)$.
- Sorting the edges in ascending order takes $O(|E| \log |E|) = O(|E| \log |V|)$ time.
- The for loop has $|E|$ iterations.
- **Key questions:**
 - How to implement an iteration of the for loop?
 - How to determine if adding an edge creates a cycle?

Implementing Kruskal's Algorithm

Observation: During the execution of the algorithm, the set of edges added to the solution (red edges) forms a set of disjoint sub-trees of the MST.

- To determine if adding an edge creates a cycle is the same as to determine if its end points are in the same sub-tree.




In this step, there are 4 disjoint sub-trees.

Implementing Kruskal's Algorithm

Observation: During the execution of the algorithm, the set of edges added to the solution (red edges) forms a set of disjoint sub-trees of the MST.

- To determine if adding an edge creates a cycle is the same as to determine if its end points are in the same sub-tree.




In this step, there are 4 disjoint sub-trees.

Implementing Kruskal's Algorithm

Observation: During the execution of the algorithm, the set of edges added to the solution (red edges) forms a set of disjoint sub-trees of the MST.

- To determine if adding an edge creates a cycle is the same as to determine if its end points are in the same sub-tree.




In this step, there are 4 disjoint sub-trees.

Implementing Kruskal's Algorithm

Observation: During the execution of the algorithm, the set of edges added to the solution (red edges) forms a set of disjoint sub-trees of the MST.

- To determine if adding an edge creates a cycle is the same as to determine if its end points are in the same sub-tree.




In this step, there are 4 disjoint sub-trees.

Implementing Kruskal's Algorithm

Observation: During the execution of the algorithm, the set of edges added to the solution (red edges) forms a set of disjoint sub-trees of the MST.

- To determine if adding an edge creates a cycle is the same as to determine if its end points are in the same sub-tree.




In this step, there are 4 disjoint sub-trees.

Implementing Kruskal's Algorithm

Idea: Design a data structure that remembers the current set of disjoint sub-trees such that we can efficiently

- 1) find the subtree that a vertex, say, c , belongs to, through a procedure **find-set**(c); and
- 2) merge two subtrees through **union**(**find-set**(c), **find-set**(d)).



Implementing Kruskal's Algorithm

- 1) **initialize** $T = \{ \}$.
- 2) **for** all vertices v : **initialize** a sub-tree for v via **make-set**(v).
- 3) **for** all edges (u, v) in ascending order of weights :
- 4) **if** **find-set**(u) \neq **find-set**(v) :
- 5) add (u, v) to T ;
- 6) **union**(**find-set**(u), **find-set**(v)).

Running Time:

- # of **make-set**: $|V|$; # of **find-set**: $2|E|$; # of **union**: $|V|-1$.
- Suppose the data structure implements these subroutine in $O(\log |V|)$ time, then the total running time is $O(|E| \log |V|)$.

A Data Structure of Disjoint Sets

We need to maintain a collection of sets from n elements (vertices). Our data structure must support the followings:

- Given any two elements x, y , we need to determine whether **find-set**(x) = **find-set**(y), i.e., to determine whether x and y belongs to the same set.
- Given any two sets in the current collection, we need to replace these two sets by its **union**.


High-Level Approach

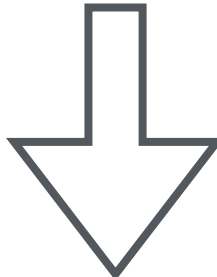
Idea: Maintain a tree for the vertices in each set and name each set after the root vertex.

- For **find-set**(x), we just trace back to the root.
- To **union** two sets, we append the root of one set to be a child of the root of the other set.

Sample Run


$\{a\}$ $\{b\}$ $\{c\}$ $\{d\}$ $\{e\}$ $\{f\}$ $\{g\}$





union(find(*a*), find(*e*)) 

Sample Run

$\{a\} \{b\} \{c\} \{d\} \{e\} \{f\} \{g\}$




union(find(*a*), find(*e*))




Sample Run

$\{a, e\}$ $\{b\}$ $\{c\}$ $\{d\}$ $\{f\}$ $\{g\}$




union(find(*d*), find(*g*))





Sample Run

$\{a, e\}$ $\{b\}$ $\{c\}$ $\{d\}$ $\{f\}$ $\{g\}$




union(find(d), find(g))




Sample Run

$\{a, e\}$ $\{b\}$ $\{c\}$ $\{d, g\}$ $\{f\}$




union(find(*e*), find(*g*))





Sample Run

{a, e} {b} {c} {d, g} {f}




union(find(e), find(g))



Summary

- To execute **find-set**(x), we traverse the parent pointers from x up to the root, and the root is used as the name of the set.
- So, **find-set**(x) = **find-set**(y) if and only if the root returned by **find-set**(x) is equal to that returned by **find-set**(y).
- To execute **union**(r_1, r_2), we make a root to be the child of the other root.




Question: Which vertex shall be the new root?

Running Time Analysis

To answer the question, we need to first understand how the running time of **find-set** and **union** depends on the structure of the trees that we maintain.


- Running time for a **union** operation: $O(1)$.
- Running time for a **find-set**(x):
the height of the tree containing x , which can be $O(n)$.



Running Time Analysis

To answer the question, we need to first understand how the running time of **find-set** and **union** depends on the structure of the trees that we maintain.


- Running time for a **union** operation: $O(1)$.
- Running time for a **find-set**(x):
the height of the tree containing x , which can be $O(n)$.



Running Time Analysis

To answer the question, we need to first understand how the running time of **find-set** and **union** depends on the structure of the trees that we maintain.

- Running time for a **union** operation: $O(1)$.
- Running time for a **find-set**(x):
the height of the tree containing x , which can be $O(n)$.




union(find-set(e), find-set(f))
union(find-set(d), find-set(e))

Running Time Analysis

To answer the question, we need to first understand how the running time of **find-set** and **union** depends on the structure of the trees that we maintain.

- Running time for a **union** operation: $O(1)$.
- Running time for a **find-set**(x):
the height of the tree containing x , which can be $O(n)$.




union(find-set(e), find-set(f))
union(find-set(d), find-set(e))
union(find-set(c), find-set(d))

Running Time Analysis

To answer the question, we need to first understand how the running time of **find-set** and **union** depends on the structure of the trees that we maintain.

- Running time for a **union** operation: $O(1)$.
- Running time for a **find-set**(x):
the height of the tree containing x , which can be $O(n)$.




union(find-set(e), find-set(f))
union(find-set(d), find-set(e))
union(find-set(c), find-set(d))
union(find-set(b), find-set(c))

Running Time Analysis

To answer the question, we need to first understand how the running time of **find-set** and **union** depends on the structure of the trees that we maintain.

- Running time for a **union** operation: $O(1)$.
- Running time for a **find-set**(x):
the height of the tree containing x , which can be $O(n)$.




union(find-set(e), find-set(f))
union(find-set(d), find-set(e))
union(find-set(c), find-set(d))
union(find-set(b), find-set(c))
union(find-set(a), find-set(b))

Running Time Analysis

To answer the question, we need to first understand how the running time of **find-set** and **union** depends on the structure of the trees that we maintain.

- Running time for a **union** operation: $O(1)$.
- Running time for a **find-set**(x):
the height of the tree containing x , which can be $O(n)$.





union(find-set(e), find-set(f))
union(find-set(d), find-set(e))
union(find-set(c), find-set(d))
union(find-set(b), find-set(c))
union(find-set(a), find-set(b))

Then, find-set(f)
takes $n-1$ steps.

Can we guarantee small height?

- If we could implement **union** such that the trees have small height, then **find-set** would have small running time.
- **How?** One way to do it is the **union-by-size** heuristic:
 - To **union**(s_1, s_2), we make the tree with smaller size the child of the one with larger size.
 - Break ties arbitrarily.
 - **Example:**




Revisiting the $O(n)$ -Time Example

**union(find-set(e), find-set(f)), union(find-set(d), find-set(e)),
union(find-set(c), find-set(d)), union(find-set(b), find-set(c)),
union(find-set(a), find-set(b))**




Revisiting the $O(n)$ -Time Example

**union(find-set(e), find-set(f)), union(find-set(d), find-set(e)),
union(find-set(c), find-set(d)), union(find-set(b), find-set(c)),
union(find-set(a), find-set(b))**




Revisiting the $O(n)$ -Time Example

`union(find-set(e), find-set(f))`, **`union(find-set(d), find-set(e))`**,
`union(find-set(c), find-set(d))`, **`union(find-set(b), find-set(c))`**,
`union(find-set(a), find-set(b))`




Revisiting the $O(n)$ -Time Example

`union(find-set(e), find-set(f)), union(find-set(d), find-set(e)),`
`union(find-set(c), find-set(d)), union(find-set(b), find-set(c)),`
`union(find-set(a), find-set(b))`




Revisiting the $O(n)$ -Time Example

`union(find-set(e), find-set(f)), union(find-set(d), find-set(e)),
union(find-set(c), find-set(d)), union(find-set(b), find-set(c)),
union(find-set(a), find-set(b))`




Revisiting the $O(n)$ -Time Example

`union(find-set(e), find-set(f)), union(find-set(d), find-set(e)),
union(find-set(c), find-set(d)), union(find-set(b), find-set(c)),
union(find-set(a), find-set(b))`



Revisiting the $O(n)$ -Time Example

`union(find-set(e), find-set(f)), union(find-set(d), find-set(e)),
union(find-set(c), find-set(d)), union(find-set(b), find-set(c)),
union(find-set(a), find-set(b))`



The union-by-size heuristic
gives a tree with height 1!

Analyzing the Union-By-Size Heuristic

Claim. If we use union-by-size, then any tree with height h must have at least 2^h nodes.


Proof (by induction).

Base case: $h = 0$ is true as any tree has at least 1 node.

Inductive hypothesis: The claim is true for trees with height h .

Inductive step: We will show it is true for trees with height $h+1$.

- **Key observation:** We get a tree with height $h+1$ only when we union two trees and one of them has height h .



Analyzing the Union-By-Size Heuristic

Claim. If we use union-by-size, then any tree with height h must have at least 2^h nodes.

Corollary. If we use union-by-size, then any tree must have height $h \leq \log n$.

Proof.

- Since there are only n nodes, the size of any tree is $\leq n$.
- Together with the above claim, we have $n \geq \text{tree size} \geq 2^h$.
- Equivalently, we have $\log n \geq \log(\text{tree size}) \geq h$.

Summary

- If we implement the Disjoint Sets data structure using the union-by-size heuristic, then
 - **find-set** runs in time $O(\log n)$;
 - **union** runs in time $O(1)$.
- Substitute these bounds in our analysis of the running time, we conclude that Kruskal's algorithm runs in $O(m \log n)$ time.